

RAM PRASATH K

ramk612000@gmail.com | (+91) 9943933752 | Pune, India | Machine Learning Engineer

[LinkedIn](#) | [GitHub](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

ML Engineer with 2+ years of production experience specializing in Agentic AI systems, optimized LLM/VLM inference pipelines, and vision-action models. Handled full-lifecycle deployment—from fine-tuning (GRPO, DPO, QLoRA) on AMD MI300X/RTX hardware to serving quantized models (vLLM, W8A8) at the edge. **1st Place Winner at the Mistral AI Worldwide Hackathon 2026**. Targeting core Agentic AI, AIOps, and Infra Intelligence roles.

EXPERIENCE

Jr. R&D Engineer - Ubisoft India Studios

Oct 2025 - Present

Game AI & Quality Control division - building autonomous gameplay agents

- Contributing to **EagleEye**, an internal clipping detection tool, by fine-tuning Qwen-VL for visual clipping classification and exploring XGBoost pipeline on manual asset data IoU, scale, and categorical asset features to improve detection of intentional vs. unintentional clipping cases.
- Researched and built a prototype fine-tuning **NVIDIA NitroGen, a diffusion-transformer Vision-Action foundation model**, to adapt to Ubisoft game environments, targeting reduction in manual QA regression testing time.
- Partnered with gameplay and QA engineers to scope requirements and deliver internal tooling including asset validation scripts for **Anvil**, changelist summarizers, and D-test case generators based on production feedback.

Stack: PyTorch · Diffusion Transformers · C# · C++ · Vision-Action Models

Data Science Engineer - Ninestars Information Technologies

Jul 2024 - Sep 2025

R&D division focused on production AI systems for OCR, document intelligence, and multi-agent automation

AI-Driven OCR & Document Understanding

- Optimized Vision-Language Model inference using vLLM with W8A16 / W4A16 quantization and served on RTX 4090; profiled tokens/sec, TTFT, and throughput trade-offs to deliver stable low-latency serving under tight GPU memory constraints.
- Developed end-to-end OCR and document understanding pipelines using Qwen-VL / Qwen2.5-VL and a custom Seq2Seq Transformer for OCR correction, NER, and structured extraction -cutting manual review by 20% and boosting entity extraction precision by 15%.
- Built a layout-aware annotation acceleration system using YOLO and RT-DETR integrated with a Flask-based backend, achieving a 4x increase in labeling throughput; partnered with the annotation team to scope requirements and iterate on delivery.

Multi-Agent Platform

- Built prototype for multi-agent automation platform on Google ADK + MCP, integrating locally-served GPT-OSS-20B, Ollama, , REST API, WebSockets, and vLLM implementing ReAct reasoning, self-reflection loops, and hierarchical patterns across the orchestration layer.
- Built evaluation pipelines and observability hooks for agentic workflows, tracking accuracy, latency, and tool-call success rates against production readiness thresholds.

Stack: Python · PyTorch · Hugging Face · vLLM · Google ADK · MCP · GPT-OSS-20B · Ollama · AWS · FastAPI · Flask · Docker · WebSockets · YOLO · RT-DETR · PostgreSQL · W&B · Langfuse

AI/ML Research Engineer (Intern) - Bosch Global Software Technologies

Aug 2023 - Jul 2024

HMI division -11-month end-to-end ownership from research to NXP edge deployment

Voice Biometrics

- Built a ResNet-SE automotive voice authentication system with anti-spoofing -achieved 95% accuracy and 4% Equal Error Rate (EER) under real-world noise conditions.
- Trained and deployed multilingual speech models (VoxCeleb, IndicSuperb datasets); exported optimized inference via ONNX / TensorFlow Lite for NXP edge hardware, and the edge deployment lifecycle.
- Worked with HMI and embedded engineering teams to scope latency and memory budgets, and shipped validated models onto target hardware.

Stack: PyTorch · Torchaudio · ResNet-SE · ONNX · TensorFlow Lite · NXP Edge Hardware

TECHNICAL SKILLS

Languages: Python, C++, C#, SQL

ML Engineering & Serving: vLLM, FastAPI, Flask, REST APIs, AWS, MCP Servers, OpenAI-compatible APIs, scikit-learn, Jupyter

Agentic Systems: Google ADK, MCP, ReAct, Self-Reflection, Multi-Agent Orchestration, Tool-Use, RAG

Model Optimization: W8A8, W4A16, AWQ, ONNX, TensorFlow, Knowledge Distillation, Quantization-Aware Training

Training & Fine-tuning: PyTorch, Hugging Face Transformers, SFT, LoRA, QLoRA, DPO, GRPO, FSDP.

Evaluation & Observability: LLM-native metrics (tokens/sec, TTFT, cost-per-request), Eval Pipelines, W&B Tracing.

Data & Storage: PostgreSQL, GraphDB, Neo4j, Vector Stores (Qdrant)

MLOps & Tooling: Weights & Biases, Git, CI/CD, Linux, Docker

ACHIEVEMENTS

- Mistral AI Worldwide Hackathon 2026 - Winner, #1 (Online Track): Built Mistral Raid, an AI-powered dungeon crawler where the boss dynamically adapts to player voice and gameplay using Mistral LLM agents, Voxtral STT, and real-time combat telemetry. Designed real-time telemetry pipeline ingesting structured data points - map position, player health, etc to Mistral LLM agents, enabling dynamic boss adaptation based on live combat state.
- AMD, Unsloth, PyTorch Synthetic Data AI Agents Challenge (Nov 2025): Designed and fine-tuned QA AI agents using synthetic data pipelines and Synthetic-Data-Kit on AMD Instinct MI300X GPUs, evaluated SFT, LoRA, QLoRA, and DPO fine-tuning workflows on next-gen GPU architecture.
- NVIDIA Nemotron Model Reasoning Challenge - Kaggle (2025–2026): Improved Nemotron 3 Nano reasoning accuracy to 85%, using prompting, synthetic data generation, data curation, and fine-tuning (SFT, LoRA, GRPO) on RTX Pro 6000.

PROJECTS

XFINITE-OCR Professional - Production Document Intelligence System

- Designed and deployed a vLLM-powered Vision-Language Model (VLM) inference system with an OpenAI-compatible API, enabling scalable multimodal document understanding under tight GPU memory constraints.
- Implemented decoding-time inference controls to improve OCR robustness; conducted inference profiling studies tracking tokens/sec, TTFT, memory usage, and throughput trade-offs across batch sizes and context lengths.

Multimodal Video Captioning System with Commonsense Reasoning

- Developed a multimodal captioning system integrating visual representations and textual priors via a hybrid ResNet152 + Transformer decoder architecture, enhancing semantic relevance and caption coherence.

Student Evaluation Framework with RAG & GraphDB

- Designed a memory-augmented evaluation system combining Retrieval-Augmented Generation (RAG) with a Graph Database, enabling structured knowledge retrieval and relational reasoning improving response consistency over flat vector only retrieval.

PUBLICATIONS & PATENT

- Appearance-Trajectory Network for Video Anomaly Detection System, ICCCNT 2024, IEEE, IIT Mandi (July 2024): Memory-augmented CNN for unsupervised video surveillance achieving 82% AUC (appearance) and 84.9% AUC (skeleton), with integrated video captioning for explainability.
- Simultaneous Voice Authentication and Decision-Making Based on Intent & Entity Classification, Patent No. 202441102999 (filed Dec 24, 2024): System combining voice authentication with real-time intent and entity classification for secure, context-aware voice-driven decision automation.

EDUCATION

Indian Institute of Technology - Madras (IIT-Madras)

Bachelor of Science in Data Science and Applications (Online / Distance programme)

Dec 2024

Amrita Vishwa Vidyapeetham, School of Artificial Intelligence

Master of Technology in Data Science - CGPA 8.5 / 10

Jun 2024

SASTRA Deemed University

Bachelor of Technology in Mechanical Engineering

Jun 2022